

MFE R Programming Workshop

Week 3

Brett Dunn

Fall 2016

Data Munging

Hadley Wickham

- ▶ **Hadley Wickham** is practically famous in the R world
- ▶ He's developed a ridiculous number of useful packages
 - ▶ e.g. `ggplot2`
- ▶ Today we will look at `dplyr` and `tidyr`

dplyr

dplyr intro

- ▶ dplyr is a package for data manipulation
- ▶ data.table is another fantastic package of this type
- ▶ I'll post a solution to today's lab using both
- ▶ These slides are a cut down version of the [dplyr introduction vignette](#)

Data: nycflights13

- ▶ To explore the basic data manipulation verbs of dplyr, we'll start with the built in 'nycflights13' data frame
- ▶ This dataset contains all flights that departed from New York City in 2013

```
library(dplyr)
library(nycflights13)
```

```
head(flights,4)
```

```
## Source: local data frame [4 x 19]
```

```
##
```

```
##   year month   day dep_time sched_dep_time dep_delay
##   (int) (int) (int)   (int)         (int)         (dbl)
## 1  2013     1     1     517           515             2
## 2  2013     1     1     533           529             4
## 3  2013     1     1     542           540             2
## 4  2013     1     1     544           545            -1
```

Single table verbs

- ▶ Dplyr aims to provide a function for each basic verb of data manipulation:
 - ▶ `filter()` (and `slice()`)
 - ▶ `arrange()`
 - ▶ `select()` (and `rename()`)
 - ▶ `distinct()`
 - ▶ `mutate()` (and `transmute()`)
 - ▶ `summarise()`
 - ▶ `sample_n()` and `sample_frac()`

Filter rows with `filter()`

- ▶ `filter()` allows you to select a subset of rows in a data frame.
- ▶ The first argument is the name of the data frame.
- ▶ The second and subsequent arguments are the expressions that filter the data frame
- ▶ Select all flights on January 1st with:

```
filter(flights, month == 1, day == 1)
```

```
## Source: local data frame [842 x 19]
```

```
##
```

##	year	month	day	dep_time	sched_dep_time	dep_delay
##	(int)	(int)	(int)	(int)	(int)	(dbl)
## 1	2013	1	1	517	515	2
## 2	2013	1	1	533	529	4
## 3	2013	1	1	542	540	2
## 4	2013	1	1	544	545	-1
## 5	2013	1	1	554	600	-6
## 6	2013	1	1	554	558	-4

Select rows by position

- To select rows by position, use `slice()`

```
slice(flights, 1:10)
```

```
## Source: local data frame [10 x 19]
```

```
##
```

```
##      year month   day dep_time sched_dep_time dep_delay
##      (int) (int) (int)   (int)         (int)         (dbl)
## 1   2013     1     1     517             515             2
## 2   2013     1     1     533             529             4
## 3   2013     1     1     542             540             2
## 4   2013     1     1     544             545            -1
## 5   2013     1     1     554             600            -6
## 6   2013     1     1     554             558            -4
## 7   2013     1     1     555             600            -5
## 8   2013     1     1     557             600            -3
## 9   2013     1     1     557             600            -3
## 10  2013     1     1     558             600            -2
```

Arrange rows with `arrange()`

- ▶ `arrange()` works similarly to `filter()` except that instead of filtering or selecting rows, it reorders them

```
arrange(flights, year, month, day)
```

```
## Source: local data frame [336,776 x 19]
```

```
##
```

```
##      year month   day dep_time sched_dep_time dep_delay
##      (int) (int) (int)   (int)           (int)         (dbl)
## 1   2013     1     1     517             515           2
## 2   2013     1     1     533             529           4
## 3   2013     1     1     542             540           2
## 4   2013     1     1     544             545          -1
## 5   2013     1     1     554             600          -6
## 6   2013     1     1     554             558          -4
## 7   2013     1     1     555             600          -5
## 8   2013     1     1     557             600          -3
## 9   2013     1     1     557             600          -3
```

Use desc() to order a column in descending order

```
arrange(flights, desc(arr_delay))
```

```
## Source: local data frame [336,776 x 19]
```

```
##
```

```
##      year month   day dep_time sched_dep_time dep_delay
```

```
##      (int) (int) (int)      (int)              (int)        (dbl)
```

```
## 1   2013     1     9      641              900         1301
```

```
## 2   2013     6    15     1432             1935         1137
```

```
## 3   2013     1    10     1121             1635         1126
```

```
## 4   2013     9    20     1139             1845         1014
```

```
## 5   2013     7    22      845             1600         1005
```

```
## 6   2013     4    10     1100             1900          960
```

```
## 7   2013     3    17     2321              810          911
```

```
## 8   2013     7    22     2257              759          898
```

```
## 9   2013    12     5      756             1700          896
```

```
## 10  2013     5     3     1133             2055          878
```

```
## .. ... .. ... .. ... .. ...
```

```
## Variables not shown: arr_time (int), sched_arr_time (int)
```

Select columns with select()

- ▶ 'select()' allows you to rapidly zoom in on a useful subset using operations that usually only work on numeric variable positions:

```
# Select columns by name
```

```
select(flights, year, month, day)
```

```
## Source: local data frame [336,776 x 3]
```

```
##
```

```
##      year month   day
```

```
##      (int) (int) (int)
```

```
## 1    2013     1     1
```

```
## 2    2013     1     1
```

```
## 3    2013     1     1
```

```
## 4    2013     1     1
```

```
## 5    2013     1     1
```

```
## 6    2013     1     1
```

```
## 7    2013     1     1
```

```
## 8    2013     1     1
```

You can rename variables with `rename()`

```
rename(flights, tail_num = tailnum)
```

```
## Source: local data frame [336,776 x 19]
```

```
##
```

```
##      year month   day dep_time sched_dep_time dep_delay
```

```
##      (int) (int) (int)      (int)              (int)      (dbl)
```

```
## 1   2013     1     1      517                515          2
```

```
## 2   2013     1     1      533                529          4
```

```
## 3   2013     1     1      542                540          2
```

```
## 4   2013     1     1      544                545         -1
```

```
## 5   2013     1     1      554                600         -6
```

```
## 6   2013     1     1      554                558         -4
```

```
## 7   2013     1     1      555                600         -5
```

```
## 8   2013     1     1      557                600         -3
```

```
## 9   2013     1     1      557                600         -3
```

```
## 10  2013     1     1      558                600         -2
```

```
## ..      ...      ...      ...      ...      ...      ...
```

```
## Variables not shown: arr_time (int), sched_arr_time (int)
```

Extract distinct (unique) rows

- ▶ A common use of 'select()' is to find the values of a set of variables.
- ▶ This is particularly useful in conjunction with the distinct() verb

```
distinct(select(flights, tailnum))
```

```
## Source: local data frame [4,044 x 1]
```

```
##
```

```
##      tailnum
```

```
##      (chr)
```

```
## 1    N14228
```

```
## 2    N24211
```

```
## 3    N619AA
```

```
## 4    N804JB
```

```
## 5    N668DN
```

```
## 6    N39463
```

```
## 7    N516JB
```

Add new columns with mutate()

```
mutate(flights,  
  gain = arr_delay - dep_delay,  
  speed = distance / air_time * 60)
```

```
## Source: local data frame [336,776 x 21]
```

```
##
```

##	year	month	day	dep_time	sched_dep_time	dep_delay
##	(int)	(int)	(int)	(int)	(int)	(dbl)
## 1	2013	1	1	517	515	2
## 2	2013	1	1	533	529	4
## 3	2013	1	1	542	540	2
## 4	2013	1	1	544	545	-1
## 5	2013	1	1	554	600	-6
## 6	2013	1	1	554	558	-4
## 7	2013	1	1	555	600	-5
## 8	2013	1	1	557	600	-3
## 9	2013	1	1	557	600	-3
## 10	2013	1	1	558	600	-2

If you only want to keep the new variables, use `transmute()`

```
transmute(flights,  
  gain = arr_delay - dep_delay,  
  gain_per_hour = gain / (air_time / 60)  
)
```

```
## Source: local data frame [336,776 x 2]
```

```
##
```

```
##      gain gain_per_hour
```

```
##      (dbl)          (dbl)
```

```
## 1         9         2.378855
```

```
## 2        16         4.229075
```

```
## 3        31        11.625000
```

```
## 4       -17        -5.573770
```

```
## 5       -19        -9.827586
```

```
## 6        16         6.400000
```

```
## 7        24         9.113924
```

```
## 8         11        12.452820
```


Summarise values with summarise()

- ▶ The last verb is 'summarise()}. It collapses a data frame to a single row:

```
summarise(flights,  
  delay = mean(dep_delay, na.rm = TRUE))
```

```
## Source: local data frame [1 x 1]  
##  
##      delay  
##      (dbl)  
## 1 12.63907
```

Commonalities

- ▶ The syntax and function of all these verbs are very similar:
 - ▶ The first argument is a data frame.
 - ▶ The subsequent arguments describe what to do with the data frame.
 - ▶ The result is a new data frame
- ▶ Together these properties make it easy to chain together multiple simple steps to achieve a complex result.

Grouped operations

- ▶ These verbs are useful on their own, but they become really powerful when you apply them to groups of observations
- ▶ In dplyr, you do this by with the `group_by()` function
- ▶ It breaks down a dataset into specified groups of rows

Grouped operations (cont.)

- ▶ Grouping affects the verbs as follows:
 - ▶ `grouped select()` is the same as `ungrouped select()`, except that grouping variables are always retained.
 - ▶ `grouped arrange()` orders first by the grouping variables
 - ▶ `mutate()` and `filter()` are most useful in conjunction with window functions (like `rank()`, or `min(x) = x`). They are described in detail in `vignette("window-functions")`.
 - ▶ `sample_n()` and `sample_frac()` sample the specified number/fraction of rows in each group.
 - ▶ `slice()` extracts rows within each group.
 - ▶ `summarise()` is powerful and easy to understand, as described in more detail below.

group_by Example

- For example, we could use these to find the number of planes and the number of flights that go to each possible destination:

```
destinations <- group_by(flights, dest)
summarise(destinations,
  planes = n_distinct(tailnum),
  flights = n()
)
```

```
## Source: local data frame [105 x 3]
```

```
##
```

```
##      dest planes flights
```

```
##      (chr)   (int)   (int)
```

```
## 1     ABQ     108     254
```

```
## 2     ACK      58     265
```

```
## 3     ALB     172     439
```

```
## 4     ANC       6       8
```

```
## 5     ATL    1180    17215
```

Chaining

- ▶ The dplyr API is functional — function calls don't have side-effects.
- ▶ You must always save their results. **UGLY**
- ▶ To get around this problem, dplyr provides the `%>%` operator
- ▶ `x %>% f(y)` turns into `f(x, y)`

```
flights %>%  
  group_by(year, month, day) %>%  
  select(arr_delay, dep_delay) %>%  
  summarise(arr = mean(arr_delay, na.rm = TRUE),  
            dep = mean(dep_delay, na.rm = TRUE)) %>%  
  filter(arr > 30 | dep > 30)
```

```
## Source: local data frame [49 x 5]
```

```
## Groups: year, month [11]
```

```
##
```

```
##   year month   day    arr    dep  
##   (int) (int) (int)  (dbl)  (dbl)
```

Multiple table verbs

- ▶ dplyr implements the four most useful SQL joins:
 - ▶ `inner_join(x, y)`: matching $x + y$
 - ▶ `left_join(x, y)`: all $x +$ matching y
 - ▶ `semi_join(x, y)`: all x with match in y
 - ▶ `anti_join(x, y)`: all x without match in y
- ▶ And provides methods for:
 - ▶ `intersect(x, y)`: all rows in both x and y
 - ▶ `union(x, y)`: rows in either x or y
 - ▶ `setdiff(x, y)`: rows in x , but not y

tidyr

Sample data

```
library(tidyr)

stocks <- data.frame(
  time = as.Date('2009-01-01') + 0:9,
  X = rnorm(10, 0, 1),
  Y = rnorm(10, 0, 2),
  Z = rnorm(10, 0, 4)
)

stocks
```

##		time	X	Y	Z
## 1		2009-01-01	1.06263264	0.3846354	0.6182198
## 2		2009-01-02	-0.75972633	4.1446913	-4.4182804
## 3		2009-01-03	0.04539981	-0.5124680	0.7407128
## 4		2009-01-04	-0.10552157	-0.9878174	-2.9890774
## 5		2009-01-05	0.53070515	0.3024455	-10.6035637
## 6		2009-01-06	-2.81381162	3.1397657	-2.5040303

Bring columns together with `gather()`

```
stocksm <- stocks %>% gather(stock, price, -time)
stocksm
```

##		time	stock	price
## 1		2009-01-01	X	1.06263264
## 2		2009-01-02	X	-0.75972633
## 3		2009-01-03	X	0.04539981
## 4		2009-01-04	X	-0.10552157
## 5		2009-01-05	X	0.53070515
## 6		2009-01-06	X	-2.81381162
## 7		2009-01-07	X	-1.94841353
## 8		2009-01-08	X	-0.98660587
## 9		2009-01-09	X	0.34193470
## 10		2009-01-10	X	-2.05957063
## 11		2009-01-01	Y	0.38463544
## 12		2009-01-02	Y	4.14469129
## 13		2009-01-03	Y	-0.51246801
## 14		2009-01-04	Y	-0.98781738

Split a column with spread()

```
stocksm %>% spread(stock, price)
```

##		time	X	Y	Z
## 1		2009-01-01	1.06263264	0.3846354	0.6182198
## 2		2009-01-02	-0.75972633	4.1446913	-4.4182804
## 3		2009-01-03	0.04539981	-0.5124680	0.7407128
## 4		2009-01-04	-0.10552157	-0.9878174	-2.9890774
## 5		2009-01-05	0.53070515	0.3024455	-10.6035637
## 6		2009-01-06	-2.81381162	3.1397657	-2.5040303
## 7		2009-01-07	-1.94841353	-1.7766495	9.8588272
## 8		2009-01-08	-0.98660587	0.1021915	-0.2715888
## 9		2009-01-09	0.34193470	3.3832790	-5.5933985
## 10		2009-01-10	-2.05957063	1.9892377	-2.1889344

```
stocksm %>% spread(time, price)
```

##	stock	2009-01-01	2009-01-02	2009-01-03	2009-01-04	2009-01-05
## 1	X	1.0626326	-0.7597263	0.04539981	-0.1055216	0.53070515

spread() and gather() are complements

```
df <- data.frame(x = c("a", "b"), y = c(3, 4), z = c(5, 6))
df
```

```
##    x y z
## 1 a 3 5
## 2 b 4 6
```

```
df %>% spread(x, y) %>% gather(x, y, a:b, na.rm = TRUE)
```

```
##    z x y
## 1 5 a 3
## 4 6 b 4
```

There's much more

- ▶ As usual, read the [vignette](#) on the CRAN page