

MFE R Programming Workshop

Week 2

Brett R. Dunn

Fall 2016

Introduction

Questions

Any questions before we start?

Overview of Week 2

- ▶ Time Series Data in R
- ▶ Retrieving Time Series Data from the Web
- ▶ Graphics in R

Time Series Data in R (xts)

What is a Time Series?

- ▶ A time series is a set of observations x_t , each one being recorded at a specified time t .

Key R Time Series Packages

- ▶ `xts`: eXtensible Time Series.
- ▶ `zoo`: Z's Ordered Observations.
 - ▶ Both were created by Achim Zeileis.
- ▶ `lubridate`
 - ▶ Created by Garrett Golemund and Hadley Wickham.

Date Classes in R

- ▶ Date is in yyyy-mm-dd format.
- ▶ POSIXct represents the (signed) number of seconds since the beginning of 1970 (in the UTC time zone) as a numeric vector.
- ▶ POSIXlt is a named list of vectors representing sec, min, hour, mday, mon, year, time zone parameters, and a few other items.

```
x <- Sys.time()  # clock time as a POSIXct object  
x
```

```
## [1] "2016-11-02 18:16:27 PDT"
```

```
class(x)
```

```
## [1] "POSIXct" "POSIXt"
```


What is xts?

- ▶ xts is an extended zoo object.
- ▶ A zoo object is a matrix with a vector of times that form an index.

```
library(xts)
# xts is a matrix plus an index
x <- matrix(1:4, nrow=2, ncol=2)
idx <- seq(as.Date("2016-10-27"), length=2, by="days")
x_xts <- xts(x, order.by = idx)
x_xts
```

```
##           [,1] [,2]
## 2016-10-27     1     3
## 2016-10-28     2     4
```

Constructing xts

- ▶ The function `xts()` gives you a few other options as well.
 - ▶ See `?xts`.
 - ▶ `unique` forces times to be unique.
 - ▶ `tzzone` sets the time zone of the series.
- ▶ The index should be of class `Date`, `POSIX`, `timeDate`, `chron`, etc.
- ▶ If the dates are not in chronological order, the `xts` constructor will automatically order the time series.
- ▶ Since `xts` is a subclass of `zoo`, `xts` gives us all the functionality of `zoo`.

Deconstructing xts

- ▶ How do we get the original index and matrix back?
 - ▶ `coredata` extracts the matrix.
 - ▶ `index` extracts the index.

```
coredata(x_xts)  # Gives us a matrix
```

```
##      [,1] [,2]  
## [1,]    1    3  
## [2,]    2    4
```

```
index(x_xts)  # Gives us a vector of dates
```

```
## [1] "2016-10-27" "2016-10-28"
```

Viewing the structure of an xts Object.

- ▶ The `str()` function compactly displays the internal structure of an R object.

```
str(x_xts)
```

```
## An 'xts' object on 2016-10-27/2016-10-28 containing:  
##   Data: int [1:2, 1:2] 1 2 3 4  
##   Indexed by objects of class: [Date] TZ: UTC  
##   xts Attributes:  
##   NULL
```

Importing and Exporting Time Series

- ▶ Importing:

1. Read data into R using one of the usual functions.

- ▶ `read.table()`, `read.xts()`, `read.zoo()`, etc.

2. `as.xts()` converts R objects to `xts`.

- ▶ Exporting:

- ▶ `write.zoo(x, "file")` for text files.

- ▶ `saveRDS(x, "file")` for future use in R.

Subsetting Time Series

- xts supports one and two-sided intervals.

```
# Load fund data  
data(edhec, package = "PerformanceAnalytics")  
edhec["2007-01/2007-02", 1] # interval
```

```
##              Convertible Arbitrage  
## 2007-01-31                0.0130  
## 2007-02-28                0.0117
```

```
head(edhec["2007-01/", 1]) # start in January 2007
```

```
##              Convertible Arbitrage  
## 2007-01-31                0.0130  
## 2007-02-28                0.0117  
## 2007-03-31                0.0060  
## 2007-04-30                0.0026  
## 2007-05-31                0.0110
```

Truncated Dates

- ▶ xts allows you to truncate dates

```
# January 2007 to March  
edhec["200701/03", 1] # interval
```

##	Convertible Arbitrage
## 2007-01-31	0.0130
## 2007-02-28	0.0117
## 2007-03-31	0.0060

Other Ways to Extract Values

- We can subset xts objects with vectors of integers, logicals, or dates.

```
edhec[c(1,2), 1]  # integers
```

```
##              Convertible Arbitrage
## 1997-01-31              0.0119
## 1997-02-28              0.0123
```

```
edhec[(index(edhec) < "1997-02-28"), 1]  # a logical vector
```

```
##              Convertible Arbitrage
## 1997-01-31              0.0119
```

```
edhec[c("1997-01-31", "1997-02-28"), 1]  # a date vector
```

```
##              Convertible Arbitrage
## 1997-01-31              0.0119
```


first() and last() Functions

- ▶ R uses `head()` and `tail()` to look at the start and end of a series.
 - ▶ i.e. “the first 3 rows” or “the last 6 rows”.
- ▶ `xts` has two functions `first()` and `last()`.
 - ▶ i.e. “the first 6 days” or “the last 6 months”

```
first(edhec[, "Convertible Arbitrage" ], "3 months")
```

##	Convertible Arbitrage
## 1997-01-31	0.0119
## 1997-02-28	0.0123
## 1997-03-31	0.0078

Math Operations

- ▶ Math operations are on the intersection of times.

```
x <- edhec["199701/02", 1]
y <- edhec["199702/03", 1]
x + y # only the intersection
```

```
##              Convertible.Arbitrage
## 1997-02-28          0.0246
```

Operations on the Union

```
x + merge(y, index(x), fill = 0)
```

```
##              Convertible.Arbitrage
## 1997-01-31                0.0119
## 1997-02-28                0.0246
```

```
x + merge(y, index(x), fill = na.locf)
```

```
##              Convertible.Arbitrage
## 1997-01-31                NA
## 1997-02-28                0.0246
```

Database Joins

- ▶ There are four main database joins: inner, outer, left and right joins.
 - ▶ inner join: intersection.
 - ▶ outer join: union.
 - ▶ left: using times from the left series.
 - ▶ right: using times from the right series.

Merging xts objects

- ▶ We can merge xts objects using the merge function.
- ▶ merge takes three arguments.
 - ▶ an arbitrary number of time series.
 - ▶ fill, which handles missing data.
 - ▶ join, the type of join we want to do.

```
colnames(x) <- "x"; colnames(y) <- "y"  
merge(x, y)
```

```
##           x      y  
## 1997-01-31 0.0119 NA  
## 1997-02-28 0.0123 0.0123  
## 1997-03-31      NA 0.0078
```

Merging xts Objects: Left and Right Joins

```
merge(x, y, join='left')
```

```
##              x      y
## 1997-01-31 0.0119    NA
## 1997-02-28 0.0123 0.0123
```

```
merge(x, y, join='right')
```

```
##              x      y
## 1997-02-28 0.0123 0.0123
## 1997-03-31      NA 0.0078
```

Missing Data

- locf: last observation carried forward

```
x <- c(1, NA, NA, 4)
idx <- seq(as.Date("2016-10-27"), length=4, by="days")
x <- xts(x, order.by = idx); colnames(x) <- "x"
cbind(x, na.locf(x), na.locf(x, fromLast = TRUE))
```

	x	x.1	x.2
## 2016-10-27	1	1	1
## 2016-10-28	NA	1	4
## 2016-10-29	NA	1	4
## 2016-10-30	4	4	4

Other NA Options

```
na.fill(x, -999)
```

```
##           x
## 2016-10-27  1
## 2016-10-28 -999
## 2016-10-29 -999
## 2016-10-30  4
```

```
na.omit(x)
```

```
##           x
## 2016-10-27  1
## 2016-10-30  4
```


Interpolate NAs

```
na.approx(x)
```

```
##           x
## 2016-10-27 1
## 2016-10-28 2
## 2016-10-29 3
## 2016-10-30 4
```

Lagging a Time Series

- ▶ `lag(x, k = 1, na.pad = TRUE)`
 - ▶ `k` is the number of lags (positive = forward and negative = backward)
 - ▶ `k` can be a vector of lags
 - ▶ `'na.pad'` pads the vector back to the original size

```
x <- na.approx(x)
cbind(x, lag(x,1), lag(x,-1))
```

```
##           x x.1 x.2
## 2016-10-27 1  NA   2
## 2016-10-28 2   1   3
## 2016-10-29 3   2   4
## 2016-10-30 4   3  NA
```

Diffferencing Series

- ▶ Differencing converts levels to changes.
- ▶ see `diff.xts` for additional function arguments.

```
x <- na.approx(x)
cbind(x, diff(x))
```

```
##           x x.1
## 2016-10-27 1  NA
## 2016-10-28 2   1
## 2016-10-29 3   1
## 2016-10-30 4   1
```

Apply over Time Periods

- ▶ `period.apply()` applies a function over time intervals.
- ▶ `endpoints` gives us the row numbers of endpoints.
- ▶ `apply.monthly`, `apply.daily`, `apply.quarterly`, etc. take care of the endpoint calculation for us.

```
edhec9701 <- edhec["1997/2001", c(1,3)]  
# determine the endpoints  
ep <- endpoints(edhec9701, "years")  
period.apply(edhec9701, INDEX=ep, FUN=mean)
```

##	Convertible	Arbitrage	Distressed	Securities
## 1997-12-31	0.01159167		0.013016667	
## 1998-12-31	0.00270000		-0.001491667	
## 1999-12-31	0.01251667		0.015225000	
## 2000-12-31	0.01377500		0.004050000	
## 2001-12-31	0.01086667		0.011525000	

do.call: A Useful R Trick

- ▶ The `do.call` function allows us to specify the name of function, either as a character or an object, and provide arguments as a list.

```
do.call(mean, args= list(1:10))
```

```
## [1] 5.5
```

```
do.call("mean", args= list(1:10))
```

```
## [1] 5.5
```

Discrete Rolling Windows

- `split`, `lapply` a function (`cumsum`, `cumprod`, `cummin`, `cummax`), and recombine.

```
edhec.yrs <- split(edhec[,1], f="years")
edhec.yrs <- lapply(edhec.yrs, cumsum)
edhec.ytd <- do.call(rbind, edhec.yrs)
edhec.ytd["200209/200303", 1]
```

##	Convertible Arbitrage
## 2002-09-30	0.0322
## 2002-10-31	0.0426
## 2002-11-30	0.0677
## 2002-12-31	0.0834
## 2003-01-31	0.0283
## 2003-02-28	0.0416
## 2003-03-31	0.0505

Continuous Rolling Windows

► `rollapply(data, width, FUN, ...)`

```
rollapply(edhec["200301/06", 1], 3, mean)
```

##	Convertible Arbitrage
## 2003-01-31	NA
## 2003-02-28	NA
## 2003-03-31	0.01683333
## 2003-04-30	0.01240000
## 2003-05-31	0.01250000
## 2003-06-30	0.00760000

Lubridate

Lubridate

- ▶ Lubridate is an R package that makes it easier to work with dates and times.
- ▶ Lubridate was created by Garrett Grolemond and Hadley Wickham.

```
# install.packages("lubridate")  
library(lubridate)
```

```
##
```

```
## Attaching package: 'lubridate'
```

```
## The following object is masked from 'package:base':
```

```
##
```

```
##      date
```

Parse a date

- ▶ Lubridate accepts lots of formats

```
ymd("20110604")
```

```
## [1] "2011-06-04"
```

```
mdy("06-04-2011")
```

```
## [1] "2011-06-04"
```

```
dmy("04/06/2011")
```

```
## [1] "2011-06-04"
```

Parse a date and time

```
ymd_hms("2011-06-04 12:00:00", tz = "Pacific/Auckland")
```

```
## [1] "2011-06-04 12:00:00 NZST"
```

Extraction

```
arrive <- ymd_hms("2011-06-04 12:00:00")  
second(arrive)
```

```
## [1] 0
```

```
second(arrive) <- 25  
arrive
```

```
## [1] "2011-06-04 12:00:25 UTC"
```

Intervals

```
arrive <- ymd_hms("2011-06-04 12:00:00")  
leave <- ymd_hms("2011-08-10 14:00:00")  
interval(arrive, leave)
```

```
## [1] 2011-06-04 12:00:00 UTC--2011-08-10 14:00:00 UTC
```

Arithmetic

```
mydate <- ymd("20130130")  
mydate + days(2)
```

```
## [1] "2013-02-01"
```

```
mydate + months(5)
```

```
## [1] "2013-06-30"
```

Arithmetic

```
mydate <- ymd("20130130")  
mydate + days(1:5)
```

```
## [1] "2013-01-31" "2013-02-01" "2013-02-02" "2013-02-03"
```

End of (next) month

```
jan31 <- ymd("2013-01-31")  
jan31 + months(1)
```

```
## [1] NA
```

```
ceiling_date(jan31, "month") - days(1)
```

```
## [1] "2013-01-31"
```

```
floor_date(jan31, "month") + months(2) - days(1)
```

```
## [1] "2013-02-28"
```


Stock Market Data in R

Data from quantmod

- ▶ With quantmod we can download stock market data into xts objects.

```
library(quantmod)
getSymbols("^GSPC", src="yahoo", from = "2008-01-01")
```


```
## [1] "GSPC"
```

```
head(GSPC,3)[, 1:4]
```

##		GSPC.Open	GSPC.High	GSPC.Low	GSPC.Close
##	2008-01-02	1467.97	1471.77	1442.07	1447.16
##	2008-01-03	1447.55	1456.80	1443.73	1447.16
##	2008-01-04	1444.01	1444.01	1411.19	1411.63

A Basic Plot

```
plot(GSPC$GSPC.Close)
```



Switch Period

- ▶ `to.period` changes the periodicity of a univariate or OHLC (open, high, low, close) object.

```
eom <- to.period(GSPC, 'months')  
head(eom, 3)
```

##		GSPC.Open	GSPC.High	GSPC.Low	GSPC.Close	GSPC
##	2008-01-31	1467.97	1471.77	1270.05	1378.55	98475
##	2008-02-29	1378.60	1396.02	1316.75	1330.63	78536
##	2008-03-31	1330.45	1359.68	1256.98	1322.70	93189
##		GSPC.Adjusted				
##	2008-01-31	1378.55				
##	2008-02-29	1330.63				
##	2008-03-31	1322.70				

Plotting in R

Motivation

One skill that isn't taught in grad school is how to make a nice chart.

- Managing Director at Citigroup

What makes a chart nice?

- ▶ The reader should look at the chart and immediately understand what data are displayed.
- ▶ This means we need:
 - ▶ A clear title.
 - ▶ Clear labels for each axis (scale and units).
 - ▶ A legend if more than one time series is displayed.
 - ▶ Different colors and line formats for different time series.
 - ▶ Grid lines.
 - ▶ Labels.

Plotting Facilities in R

- ▶ R has excellent plotting methods built-in.
- ▶ I will focus on base R.
- ▶ As a next step, I recommend learning ggplot2, an excellent plotting package.
- ▶ <http://www.r-graph-gallery.com/>

Basic Plotting


- ▶ `example(plot)`
- ▶ `example(hist)`
- ▶ `?par`
- ▶ `?plot.default`

The `plot()` Function

- ▶ `plot()` is generic function, i.e. a placeholder for a family of functions.
 - ▶ the function that is actually called depends on the class of the object on which it is called.
- ▶ `plot()` works in stages.
 - ▶ you can build up a graph in stages by issuing a series of commands.
- ▶ We will see how this works with an example.

A Basic Plot

```
x <- seq(1:12)
y <- c(69, 68.7, 70.7, 73.2, 74.9, 78.2,
      82.6, 84.4, 83.5, 79, 73.6, 67.8)
plot(x, y)
```



`xlim()` and `ylim()`

Graphical paramaters


- ▶ Graphical parameters can be set as arguments to the `par` function, or they can be passed to the `plot` function.
- ▶ Make sure to read through `?par`.
- ▶ Some useful parameters:
 - ▶ `cex`: sizing of text and symbols
 - ▶ `pch`: point type.
 - ▶ `lty`: line type.
 - ▶ 0=blank, 1=solid (default), 2=dashed, 3=dotted, 4=dotdash, 5=longdash, 6=twodash
 - ▶ `lwd`: line width.
 - ▶ `mar`: margins.

pch

- pch sets how points are displayed




```
plot(x,y, pch = 16, col='darkblue')
```



Colors in R

- ▶ `colors()` returns all available color names.
- ▶ `rainbow(n)`, `heat.colors(n)`, `terrain.colors(n)` and `cm.colors(n)` return a vector of `n` contiguous colors.


```
plot(x, y, pch = 21, col=heat.colors(12),  
     cex = 2, bg = rev(heat.colors(12)))
```



lines()

- ▶ `lines()` takes coordinates and joins the corresponding points with line segments.
 - ▶ Notice, by calling `lines` after `plot` the line is on top of the points.
 - ▶ This is why we want to build the plot in stages.


```
plot(x,y, pch = 16, col='darkblue', cex=2)  
lines(x, y, col='darkgrey', lwd = 3)
```



points()

- `points` is a generic function to draw a sequence of points at the specified coordinates. The specified character(s) are plotted, centered at the coordinates.


```
plot(x,y, pch = 16, col='darkblue', cex=2)  
lines(x, y, col='darkgrey', lwd = 3)  
points(x, y, col=rainbow(12), pch=1:12, cex=3, lwd=2)
```



grid()

- ▶ grid adds a rectangular grid to an existing plot.
- ▶ ?grid for more details.


```
plot(x,y, pch = 16, col='darkblue', cex=2)  
lines(x, y, col='darkgrey', lwd = 3)  
points(x, y, col=rainbow(12), pch=1:12, cex=3, lwd=2)  
grid(col="blue", lwd=2)
```



abline()

- ▶ `abline` adds one or more straight lines through the current plot.


```
x2 <- 1:10; y2 <- 1 + 2*x2 + rnorm(10)
plot(x2,y2, pch = 16, col='darkblue')
model <- lm(y2 ~ x2)
abline(model, col="darkgrey", lwd=2)
abline(v = 5, col = "red", lty = 2)
```



Adding a Title in Lables

- ▶ Use the main argument for a title.
- ▶ Use the xlab and ylab for axis labels.


```
plot(x,y, pch = 16, col='darkblue',  
     xlab = "Month", ylab = "Temperature (degrees F)",  
     main = "Average High Temperature in Los Angeles, CA")
```



Adding a Legend: The legend() Function

- ▶ see ?legend and example(legend)


```
plot(x,y, pch = 16, col='darkblue')  
legend("topleft", inset=.01, "Average High Temperature",  
      col = "darkblue", pch = 16, bg="white", box.col="white")
```



text() and locator()

- ▶ Use the `text()` function to add text anywhere in the current graph.
- ▶ `locator()` allows you to click on a point in the chart and returns the location.


```
plot(x,y, pch = 16, col='darkblue')  
text(2,75, "ABC")
```



curve()

- ▶ With `curve()`, you can plot a function.

```
curve(cos(x), -2*pi, 2*pi)
```



Saving a Plot to a File

1. Open a file: `pdf("name.pdf")`
 2. Create the plot.
 3. Close the device with `dev.off()`
-
- ▶ You can use `dev.copy()` to save the displayed graph.
 - ▶ See `library(help = "grDevices")` for more information.

An Example of Plotting in R

- ▶ Let's plot the cumulative (gross) return of IBM and the S&P 500 since 1980.

```
library(quantmod)
getSymbols(c("^GSPC", "IBM"), src="yahoo", from = "1979-12-
```

```
## [1] "GSPC" "IBM"
```

```
adj_close <- merge(GSPC$GSPC.Adjusted, IBM$IBM.Adjusted)
daily_returns <- diff(adj_close)/lag(adj_close)
cum_ret <- cumprod(1+daily_returns[-1,])
ret1 <- xts(matrix(1, ncol=2), as.Date("1979-12-31"))
cum_ret <- (rbind(cum_ret, ret1) - 1)*100
colnames(cum_ret) <- c("GSPC", "IBM")
```

The Data


```
head(cum_ret, 9)
```

##		GSPC	IBM
##	1979-12-31	0.0000000	0.000000
##	1980-01-02	-2.0196405	-2.912622
##	1980-01-03	-2.5199194	-1.359235
##	1980-01-04	-1.3155503	-1.553402
##	1980-01-07	-1.0468816	-1.941753
##	1980-01-08	0.9357004	4.660191
##	1980-01-09	1.0283500	1.553387
##	1980-01-10	1.8065564	4.854359
##	1980-01-11	1.8343487	4.077673

Start with a Blank Chart and Build it Up

```
plot(cum_ret$IBM, xlab="", ylab = "Cumulative Net Return (I  
      main="", major.ticks="years", minor.ticks=F,  
      type="n", major.format = "%Y", auto.grid=F,  
      ylim = c(-500, 3000))  
abline(h=seq(-500,3000,500), col="darkgrey", lty=2)  
lines(cum_ret$GSPC, col="black", lwd=2)  
lines(cum_ret$IBM, col="blue", lwd=2)  
legend("topleft", inset=.02,  
      c("IBM","GSPC"), col=c("blue", "black"),  
      lwd=c(2,2),bg="white", box.col="white")
```

The Chart



Lab 2

Let's work on Lab 2.